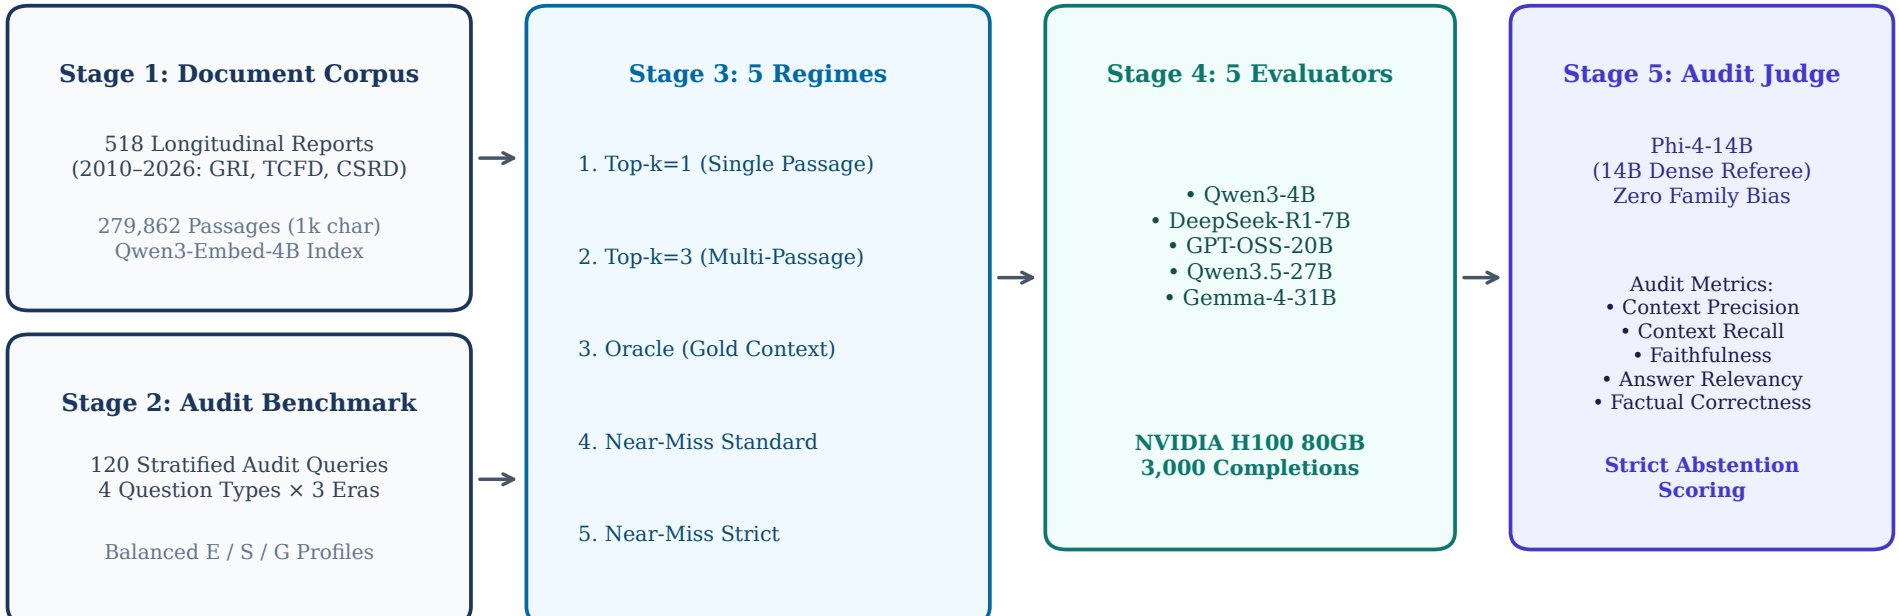
Stage 1: Document Corpus
518 Longitudinal Reports
(2010–2026: GRI, TCFD, CSRD)
279,862 Passages (1k char)
Qwen3-Embed-4B Index
Stage 2: Audit Benchmark
120 Stratified Audit Queries
4 Question Types × 3 Eras
Balanced E / S / G Profiles
Stage 3: 5 Regimes
1. Top-k=1 (Single Passage)
2. Top-k=3 (Multi-Passage)
3. Oracle (Gold Context)
4. Near-Miss Standard
5. Near-Miss Strict
Stage 4: 5 Evaluators
• Qwen3-4B
• DeepSeek-R1-7B
• GPT-OSS-20B
• Qwen3.5-27B
• Gemma-4-31B
NVIDIA H100 80GB
3,000 Completions
Stage 5: Audit Judge
Phi-4-14B
(14B Dense Referee)
Zero Family Bias
Audit Metrics:
• Context Precision
• Context Recall
• Faithfulness
• Answer Relevancy
• Factual Correctness
Strict Abstention
Scoring